

Tumor Detection With Transfer Learning

Oliver Shetler

In this article, I show how to take a general purpose image classification model (Xception in this instance), and train it to specialize in a specific task (detecting tumors in brain scans).

I obtained a data-set from Guillaume Fradet's GitHub repository. The data included only two classes: brain-scans with tumors and brain-scans without. A typical image of a brain with a tumor looks like this:




Figure 1: Tumor Example

A typical image of a healthy brain looks like this:




Figure 2: Healthy Example

Getting Started

We start the project by clearing our session and setting random seeds.

```
from tensorflow.keras.backend import clear_session
from tensorflow.random import set_seed
from numpy.random import seed
```

```
clear_session()
set_seed(4006) # comment out to use random seed
seed(4006) # comment out to use random seed
```

Next, we create an augmented data set that uses the the ImageDataGenerator function to generate images that are randomly rotated, cropped and resized to varying degrees (the training set is subjected to greater modification than the validation set, and the test set is unmodified, so that it is a fair assessment of real performance).

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.applications.vgg19 import preprocess_input

datagen_train = ImageDataGenerator(rotation_range=10,
                                   width_shift_range=0.1,
                                   height_shift_range=0.1,
                                   zoom_range=0.2,
                                   preprocessing_function=preprocess_input)

datagen_val = ImageDataGenerator(rotation_range=5,
                                 zoom_range=0.2,
                                 preprocessing_function=preprocess_input)

datagen_test = ImageDataGenerator(preprocessing_function=preprocess_input)

train_gen = datagen_train.flow_from_directory('./data/train',
                                              target_size=(224,224),
                                              color_mode='rgb',
                                              batch_size=20,
                                              class_mode='categorical',
                                              shuffle=True)

val_gen = datagen_val.flow_from_directory('./data/validate',
                                          target_size=(224,224),
                                          color_mode='rgb',
                                          batch_size=10,
                                          class_mode='categorical',
                                          shuffle=True)
```

```
test_gen = datagen_test.flow_from_directory('./data/test',
                                            target_size=(224,224),
                                            color_mode='rgb',
                                            batch_size=5,
                                            class_mode='categorical',
                                            shuffle=True)
```

```
n_classes = train_gen.num_classes
```

```
train_gen.n
```

```
Found 190 images belonging to 2 classes.
Found 40 images belonging to 2 classes.
Found 20 images belonging to 2 classes.
```

Transfer Learning

First, we create creating a base model from the Xception architecture, with its top layer removed. In the Keras API, you can remove the top layer of any architecture by setting the variable `include_top` to `True`, as see below in # [1]. Then we add a pooling layer, a dropout layer (20%) and a dense output layer using the Keras functional API, which allows you to define layers as functions, which take other layers as inputs. We then freeze the model and early stopping / model checkpoints, which save the model and stop when training no longer appreciably improves the model.

If you want to remove more than the top layer, you can look at the structure of the model using `<model_name>.summary` and then you can use a for loop to make a list of all the layers in the model up to your stopping point, and then you can put the list into the `keras.models.sequential` function to obtain a model. From there, you can proceed as shown in the cell below.

```
# define a model #
```

```
from tensorflow.keras.applications.xception import Xception
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense, Dropout
from tensorflow.keras import Model
class MCDropout(Dropout):
    def call(self, inputs):
        return super().call(inputs, training=True)
```

```

# take Xception without its output layer as the base model
base_model = Xception(weights="imagenet", include_top=False) # [1]
# add a global average pooling layer
avg = GlobalAveragePooling2D()(base_model.output)
# next, adding an MCDropout layer before the output layer
# to reduce risk of over-fitting.
dropout = MCDropout(rate=0.2)(avg)
# dropout is not standard, remove if breaks
dense = Dense(2*n_classes, activation="relu")(dropout)
output = Dense(n_classes, activation="softmax")(dense)
# change dropout to avg if broken
model = Model(inputs=base_model.input, outputs=output)

# freeze the model #

for layer in base_model.layers:
    # freeze layers in base_model
    layer.trainable = False

# set up model saving and early stopping #

from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint

early_stopping_cb = EarlyStopping(patience=20)
model_checkpoint_cb = ModelCheckpoint("model.h5", save_best_only=True)
cb_list = [early_stopping_cb, model_checkpoint_cb]

```

Next, we train the top layers only, with Xception frozen, using a learning rate of 0.2.

```

from tensorflow.keras.optimizers import SGD
# import pandas as pd

optimizer_top = SGD(lr=0.2, momentum=0.9, decay=0.01)
model.compile(loss="binary_crossentropy",
              optimizer=optimizer_top,
              metrics=["accuracy"])
history_top = model.fit(train_gen,
                       epochs=20,
                       validation_data=val_gen,
                       callbacks=cb_list)

```

```

Epoch 1/20
10/10 [=====] - 14s 1s/step
- loss: 237.6979
- accuracy: 0.5368
- val_loss: 427.3730
- val_accuracy: 0.6000
...
Epoch 19/20
10/10 [=====] - 17s 2s/step
- loss: 83.7523
- accuracy: 0.7316
- val_loss: 52.5551
- val_accuracy: 0.8000
Epoch 20/20
10/10 [=====] - 16s 2s/step
- loss: 68.1259
- accuracy: 0.7105
- val_loss: 64.5771
- val_accuracy: 0.7500

```

Finally, we train the entire model (with Xception unfrozen) using a much smaller learning rate (0.0001).

```

# train whole model with safe learning rate
# (so as not to overwhelm the pretrained weights)

# Tried lr=0.01 and decay=0.001, but error went to NAN levels.
# Tried lr=0.005 and decay=0.0005.
# Tried lr=0.001 and decay=0.0001, but loss still went to NAN.
# Switched from VGG19 to Xception model, and everything worked out well at 0.001.

for layer in base_model.layers: # for each layer in the model
    layer.trainable = True # unfreeze it

optimizer_safe = SGD(lr=0.001, momentum=0.9, decay=0.0001)
model.compile(loss="binary_crossentropy", optimizer=optimizer_safe,
              metrics=["accuracy"])
history_safe = model.fit(train_gen,
                        epochs=10, #15
                        validation_data=val_gen,
                        callbacks=cb_list)

```

```
Epoch 1/10
10/10 [=====] - 65s 6s/step
- loss: 2.1301
- accuracy: 0.7263
- val_loss: 5.0712
- val_accuracy: 0.7250
Epoch 2/10
...
Epoch 9/10
10/10 [=====] - 61s 6s/step
- loss: 0.1291
- accuracy: 0.9421
- val_loss: 0.2929
- val_accuracy: 0.8750
Epoch 10/10
10/10 [=====] - 62s 6s/step
- loss: 0.0547
- accuracy: 0.9842
- val_loss: 0.1893
- val_accuracy: 0.9750
```

The performance of the model, as seen below

```
# Evaluate the model on the test set
```

```
evaluation = model.evaluate(test_gen)
```

```
4/4 [=====]
- 1s 227ms/step
- loss: 0.6421
- accuracy: 0.9000
```

is fairly strong, with 90% accuracy.

Plotting Our Results

We plot training graphs for the loss and accuracy of the model in its two training phases.

```
# Plot the training and validation loss
```

```

from matplotlib import pyplot as plt

figure, axis = plt.subplots(1, 2, figsize=(15, 5))

figure.suptitle(" Training and Validation Loss ", fontsize=20)

axis[0].xaxis.set_major_locator(MaxNLocator(integer=True))
axis[0].plot(history_top.history["loss"], label="training loss", color="orange")
axis[0].plot(history_top.history["val_loss"], label="validation loss", color="green")
axis[0].set_xlabel("Epoch")
axis[0].set_ylabel("Loss")
axis[0].legend()
axis[0].set_title(" Training On Top Layers (Xception Frozen) ")

axis[1].xaxis.set_major_locator(MaxNLocator(integer=True))
axis[1].plot(history_safe.history["loss"], label="training loss", color="orange")
axis[1].plot(history_safe.history["val_loss"], label="validation loss", color="green")
axis[1].set_xlabel("Epoch")
axis[1].set_ylabel("Loss")
axis[1].legend()
axis[1].set_title(" Training On All Layers (Xception Unfrozen) ")

figure.show()

plt.savefig("loss.png")

```



Figure 3: Loss Plots

Plot the training and validation loss

```

from matplotlib import pyplot as plt
from matplotlib.ticker import MaxNLocator

```

```

figure, axis = plt.subplots(1, 2, figsize=(15, 5))

figure.suptitle(" Training and Validation Loss ", fontsize=20)

axis[0].xaxis.set_major_locator(MaxNLocator(integer=True))
axis[0].plot(history_top.history["accuracy"], label="training accuracy", color="orange")
axis[0].plot(history_top.history["val_accuracy"], label="validation accuracy", color="green")
axis[0].set_xlabel("Epoch")
axis[0].set_ylabel("Accuracy")
axis[0].legend()
axis[0].set_title(" Training On Top Layers (Xception Frozen) ")

axis[1].xaxis.set_major_locator(MaxNLocator(integer=True))
axis[1].plot(history_safe.history["accuracy"], label="training accuracy", color="orange")
axis[1].plot(history_safe.history["val_accuracy"], label="validation accuracy", color="green")
axis[1].set_xlabel("Epoch")
axis[1].set_ylabel("Accuracy")
axis[1].legend()
axis[1].set_title(" Training On All Layers (Xception Unfrozen) ")

figure.show()

plt.savefig("accuracy.png")

```



Figure 4: Accuracy Plots

Our transfer learning procedure has repurposed the Xception model architecture for diagnosing brain tumors, and we obtained a 90% out of sample accuracy score. Clearly, this model would not be viable as a substitute for doctor expertise. However, these kinds of models can serve as an aid to overworked professionals, and, with more resources, can often achieve very high levels of

accuracy and recall.